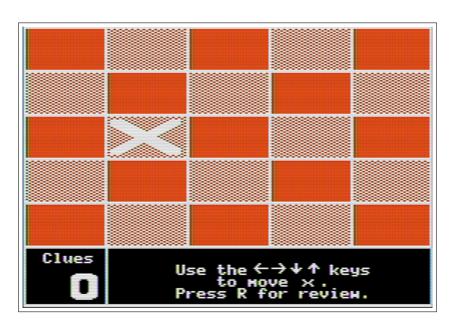
## Safari Search



2014-11-11



```
-------Safari Search-------
A 4am crack
                             2014-11-11
"Safari Search: Problem Solving and
Inference" is a 1985 educational game
distributed by Sunburst Communications.
It was designed by Thomas C. O'Brien,
with graphics and programming by Jim
Thomas.
EThe copy protection is similar to
"Exploring Science: Temperature," also
distributed by Sunburst. This write-up
is therefore quite similar to that one,
with a few corrections and the whole
thing about the disk volume number. I
COPYA fails miserably and immediately.
EDD 4 bit copy gives read errors on tracks $13 and up. Oddly, when I look
at the original disk with the Copy JC+
nibble editor, those "unreadable"
tracks do appear to be formatted and to
contain data. Needless to say, the copy
does not work.
Further inspection in the nibble editor
reveals that each track has a different
address and data prologue:
Track | Address | Data
$00 | D5 AA 96 | D5 AA AD (normal)

$01 | D5 AA 97 | D5 AA AE

$02 | D5 AA 9A | D5 AA AF

$03 | D5 AA 9B | D5 AA B2
$00
$04 | D5 AA 9D | D5 AA B3
And so on.
```

```
[S6,D1=original disk]
ES5,D1=my work disk3
JPR#5
CAPTURING BOOTØ
...reboots slot 6...
...reboots slot 5...
SAVING BOOTØ
CAPTURING BOOT1
...reboots slot 6...
...reboots slot 5...
SAVING BOOT1
SAVING RWTS
For those of you just tuning in, my
work disk runs a program I call
"AUTOTRACE" to automate the process of
boot tracing. For disks that use an
entirely custom boot process, AUTOTRACE
just captures track 0, sector 0 (saved
in a file called "BOOTO") and stops.
For "DOS 3.3-shaped" disks, which load
in the more-or-less the same way as an
unprotected DOS 3.3 disk loads, it can
also capture the next stage of the boot
process (to a file called "BOOT1").
"BOOT1" is usually sectors 0-9 on track
0, which are usually loaded into memory
at $B600..$BFFF. (Of course, there are
exceptions to every rule.)
```

Time for boot tracing with AUTOTRACE.

program (which spot-checks a few locations in memory to guess at its "normalcy"), so it extracted the RWTS routines from \$B800..\$BFFF and saved them into a third file called "RWTS". If anything looks fishy or nonstandard, ÄUTOTRACE just stops, and I have to check the files it saved so far to determine why. But in this case, it ran all the way through, automatically capturing BOOTØ, BOOTÎ, and RWTS files. Now I can use Advanced Demuffin to convert the disk to a standard format. ES6,D1=original disk3 [S6,D2=blank disk] ES5,D1=my work disk] JPR#5 Safari Search Problem Solving and Inference Designer: Thomas C. O'Brien

Graphics and programming:
\_\_\_\_\_ Jim Thomas

Press RETURN to continue.

If the boot1 code is "DOS 3.3-shaped,"
there's a good chance I'll be able to
use a tool called Advanced Demuffin to
convert the disk from whatever weird
format it uses to store its data into a
standard disk readable by unprotected
DOS 3.3. In this case, the RWTS was

close enough to normal for my AUTOTRACE

Epress "5" to switch to slot 5]

Epress 5 to switch to side 51

Epress "R" to load a new RWTS module]
 --> At \$B8, load "RWTS" from drive 1

Epress "6" to switch to slot 6]

**□**press "C" to convert disk**]** 

₃BRUN ADVANCED DEMUFFIN 1.5

This disk is 16 sectors, and the default options (copy the entire disk, all tracks, all sectors) don't need to be changed unless something goes horribly wrong.



```
+ 5:
  0123456789ABCDEF0123456789ABCDEF012
SC0:
  SC1:
  SC2:
  SC3:
  SC4:
  SC5:
  SC6:
  SC7:
  SC8:
  SC9:
  SCA:
  SCB:
  SCC:
  SCD:
  16SC $00,$00-$22,$0F BY1.0 S6,D1->S6,D2
Grr. Track $11 is generally the VTOC
(a.k.a. disk catalog), and this disk
has one sector that even the original
RWTS can't read. Either I'm incrédiblu
unlucky and got a bad original disk, ör
there's a secondary protection (nibble
check) that reads that sector in a
further non-standard way, or it's
simply unused and is intentionally bad
to make my life more difficult. The
original disk has no problem booting or
loading, so I'm quessing it's either a
nibble check or an intentional glitch.
```

The disk's own RWTS gave no read errors on any other track. This is the power and the genius of Advanced Demuffin. Every disk must be able to read itself. So, let it read itself, then capture the data and write it out in a standard format. Rebooting my work disk, I can now see the catalog on the demuffin'd copy. JPR#5 JCATALOG,86,D2 C1983 DSR^C#254 018 FREE \*A 010 LOGO \*A 009 BOOT \*A 019 CATS \*A 016 DONKEY ≭A 014 DRAGON ≭A 013 FLAMINGO ≭A 010 IGUANA \*A 019 KANGAROOS \*A 016 KITTENS \*A 015 LLAMAS \*A 014 LOON \*A 017 MENU \*A 016 RHINOS \*A 012 SEAL \*A 015 SNAILS ≭B 002 ARROWS2 ≭B 007 CAT.PICS \*B 009 DIRECTIONS.IMG \*B 009 DONKEY.PICS \*B 011 DOOR NUMBERS.IMG \*B 008 DRAGON.PICS C . . . J

```
*B
   002 ERROR STACK FIX
*B
   006 FLAMINGO.PICS
∦В
   004 HOT WARM
                 COLD
∦В
   008 IGUANA.PICS
∦В
   010 KANGAROO.PICS
   006 KITTEN.PICS
∦В
∦В
   006 LLAMA.PICS
#В
   008 LOON.PICS
#В
  002 NOT HERE
*B
   007 PACK.OBJ
∦В
   002
      PAT1
*B
   002 PAT2
∦В
   002
      PAT3
   002
∦В
       PAT4
∦В
   002
      PAT5
∦В
   002
      PAT6
∦В
   002 PAT7
∦В
   002
      PAT8
*B
  010 RHINO.PICS
*B
   010 SAFARI.DATA1
  009 SAFARI.DATA2
∦В
*B
   012 SAFARI.DATA3
   032 SAFARI.OBJ
∦В
   009 SEAL.PICS
#В
  007 SNAIL.PICS
033 SUNBURST LOGO.PIC
#В
#В
  013 TITLE.PICS
#В
*B 003 YES
            OR NO
*T 002 SET MEMLOC
Т
   002
       SOUND CHECK
This is encouraging, because it means
that T11,S0F (the unreadable sector) is
not used by the disk catalog.
No HELLO program, but that first file,
LOGO, looks promisina.
```

IRUN LOGO OK, that does work. Good. That tells me that there aren't any callbacks in the program itself that are looking for a custom DOS. (I've seen several titles like this from other publishers. They PEEK or CALL specific locations that are intentionally different from standard DOS and reboot if things look amiss.) I've cracked a number of Sunburst disks like this one. I've always gotten to this stage -- having converted the disk and confirmed that it works under standard DOS 3.3 -- then replaced tracks 0-2 with standard DOS 3.3 and declared victory. But I've never dug into the custom RWTS itself to see how it works. For example, that table of non-standard address and data prologues, where every single track uses different parameters? That's unique to Sunburst. I've never seen anything like it anywhere else. Something somewhere is changing the RWTS on the fly with every disk read. Anyway, partly because I'm curious and partly because I wanted to test my new . Post-Demuffin Patcher utility (more on that in a minute), I decided to dig into the custom RWTS on this disk. Is there a small set of changes I could make to get my demuffin'd copy to run? Or do I really need to just wipe the entire DOS and replace it?

```
3BLOAD BOOT1,A$2600
3CALL -151
*FE89G FE93G
                 ; disconnect DOS
*B600<2600.2FFFM ; move RWTS into place
The RWTS is certainly shaped like a
standard DOS 3.3 RWTS. Here's the code
at $B94F,
          which is identical to the
code I would find there on a
                               freshlu
formatted DOS 3.3. (This is where it
looks for the address proloque.)
*B94FL
B94F-
        BD 8C
                     LDA
                            $008C,X
              CØ
B952-
                     BPL
                            $B94F
        10 FB
B954-
        C9 D5
                     CMP
                            #$D5
B956-
           FЙ
                            $B948
        DØ
                     BNE
B958-
        EΑ
                     NOP
B959-
        BD 8C
               CØ.
                     LDA
                            $C08C,X
B95C-
                     BPL
        10 FB
                            $B959
B95E-
        C9 AA
                     CMP
                            #$AA
           F2
B960-
        DØ
                     BNE
                            $B954
B962-
        A0 03
                     LDY
                            #$03
B964-
        BD 8C
               CØ.
                     LDA
                            $C08C,X
B967-
        10 FB
                     BPL
                            $B964
        C9 96
B969-
                     CMP
                            #$96
           E7
                            $R954
B96B-
        DØ
                     BNE
```

**J**PR#5

```
But as I saw in the nibble editor, each
track uses a different value for the
third byte in the address prologue.
Track 0 is the normal "D5 AA 96", but
track 1 uses "D5 AA 97", track 2 uses
"D5 AA 9A", &c. The code listed above
will only read track 0. Something must
be changing the value at $B96A before
this code is run.
Turning to my trusty Disk Fixer sector
editor, I search for the hex sequence
"6A B9" and find one match on T00,804.
That sector is loaded at $BA00.
Turning back to the extracted boot1
code on my work disk, there appears to
be a routine that starts at $BA69.
*BA69L
; save accumulator
BA69- 48
                      PHA
; phase number (track number x2)
BA6A- A5 2A LDA
; divided by 2, so now just a track
; number
BA6C- 4A
                      LSR
BA6D− A8 TAY
BA6E− B9 29 BA LDA $BA29,Y
; Aha! This is modifying the third byte
; of the address prologue, just as I suspected. And it's using the track
; number as an index into a table of
; some sort.
BA71− 8D 6A B9 STA $B96A
```

```
; also changes the third byte of the
; address prologue on write
BA74- 8D 84 BC STA $BC84
; another table lookup
, anoone, dable lookup
BA77- B9 34 BA LDA $BA34,Y
; changing the third byte of the data
; prologue (read)
BA7A- 8D FC B8 STA $B8FC
; ...and write
BA7D- 8D 5D B8 STA $B85D
But wait! There's more!
; if track is $11...
                     CPY #$11
BA80- C0 11
BA82- D0 03
                    BNE $BA87
; then sector $0E is really sector $02
BA84- A9 02 LDA #$02
; (dummy opcode to hide next 2 bytes)
BA86- ĀC
; otherwise sector $0E is itself
BA87- A9 ØE LDA #$ØE
BA89- 8D CØ BF STA $BFCØ
```

WTF.

In case you have no idea what's going on here (I've never seen anything like it)... The table from \$BFB8..\$BFC7 maps physical to logical sectors. (There's an identical table on T00,S00, but it's only used by boot0 for loading boot1. This one is used by the RWTS for the lifetime of the disk.) Anyway, this code is changing that table based on the track number. So, on track \$11 (and only track \$11), physical sector \$08 is treated as sector \$02. On all other tracks, physical sector \$08 is treated as sector \$0E. No, I don't know why, except f--- you, that's why. At any rate, this confirms my theory that there's some track-based lookup table being used to modify the expected address and data prologues on each track. Fun(\*) fact: the table at \$BA29 is part of the denibblization process that turns raw nibbles on disk into values in memory, and vice-versa. So they didn't need to "waste" 35 bytes to store custom address prologues for each track; they re-used a table that the RWTS uses for an entirely different purpose -- and one that comprises valid raw nibbles. (\*) not guaranteed, actual fun may vary

```
Turning back to Disk Fixer, I searched
for "20 69 BA" (JSR $BA69) and found
nothing. But searching for "4C 69 BA"
(JMP $BA69) got two matches: one on
T00,S00 and one on T00,S08. The one on
T00,808 looks more promising. That
sector is loaded at $BE00.
"Beneath Apple DOS" (p. 8-39) says that
the routine at $BE5A is "MYSEEK", and
that it is called immediately before
the RWTS moves the drive head to the
appropriate track in preparation for a
read or write. Furthermore, my sector
search found the JMP instruction at
$BE8B, which (according to "Beneath
Apple DOS") is the final instruction of
the "MYSEEK" routine.
Looking at a freshly formatted DOS 3.3
disk, it appears that the proper JMP is
to $B9A0, which (not coincidentally) is
the final instruction of the custom
routine at $BA69, after it restores all
registers and flags:
BA8C-
        68
                    PLA.
BA8D-
        69
           ии –
                    ADC:
                          #$00
BA8F-
       48
                    PHA
BA90-
       AD 78 04
                    LDA
                          $0478
BA93-
       90
           2B
                    BCC.
                          ≴RAC0
      C9 22
BAC0-
                    CMP
                          #$22
                    ADC
BAC2- 69 00
                          #$00
BAC4-
       8D
          78 04
                    STA
                          $0478
       68
BAC7-
                    PLA
BAC8- 4C
          A0 B9
                    JMP |
                          $B9A0
```

```
This patch should restore order to the
universe:
T00,S08,$8C change 69BA to A0B9
But wait, there's more! I also spotted
this difference, in the beginning of
the routine to write the address field:
*BC69L
BC69-
           B8 B6
                    JMP
                           $B6B8
      40
BC6C-
        EΑ
                    NOP
BC6D-
       EΑ
                    NOP:
BC6E-
        EΑ
                    NOP
BC6F-
       9D 8D C0
                    STA
                           $C08D,X
           8C C0
BC72-
       DD
                    CMP
                           $008C,X
BC75-
        EΑ
                    NOP.
BC76-
        88
                    DEY
BC77-
        D0 F0
                    BNE
                           $BC69
BČ79-
       A9 D5
                    LDA
                           #$D5
                           $BCD5
BC7B-
       20 D5
              BC
                    JSR
BC7E-
      A9 AA
                    LDA
                           #$AA
       20 D5
              ВC
BC80-
                     JSR
                           $BCD5
BC83-
       A9 96
                    LDA
                           #$96
BC85-
        20
           D5
                     JSR
                           $BCD5
              BC
On a normal DOS 3.3 disk, the code at
$BC69 and $BC6C looks like this:
BC69-
       20 C3 BC
                    JSR
                           $BCC3
BC6C- 20
           C3 BC
                           $BCC3
                    JSR
```

```
$BCC3 is just an "RTS" instruction. The
entire thing is a precise wait loop to
write out self-sync bytes ($FF) before
the address field. JSR always takes 6
cycles, and RTS also always takes 6
cycles. So the two instructions wait a
total of 24 cycles before the "STA
$C08D,X" instruction writes the sync
bute to disk.
So what's at $B6B8? I'm glad you asked.
*B6B8L
                                 cycles
B6B8-
        CØ.
          05
                    CPY
                           #$05
B6BA-
        Α9
                    LDA
                           #$00
           00
                                   2(*)
2
2
3
2
3
3
3
B6BC-
           96
                    BCS
                           $B6C4
        В0
B6BE-
        EΑ
                    NOP
B6BF-
       A9 FF
                    LDA
                           #$FF
      C5
38
B6C1-
          αа
                    CMP
                           ≴00
B6C3-
                    SEC
                           ;
$B6C6 ;
B6C4- B0
                    BCS
           00
B606- 40
           6F BC
                    JMP -
                           $BC6F
(*) The Y register is 6 the first time
through this loop, so the branch at
$B6BC will be taken the first two
times, then not taken afterwards. When
the branch is taken, it requires 3 CPU
cycles and continues at $B6C4, so the
tōtal time is 2+2+3+3+3=13 cycles. When
that branch is not taken, it requires
2 CPU cycles, falls through to $B6BE,
and the whole thing takes
2+2+2+2+2+3+2+3+3=21 cycles. Adding 3
for the initial JMP at $BC69, the
whole thing (starting at $BC69, up to
but not including the STA at $BC6F)
burns either 16 ör 24 cycles.
```

```
Low-level disk activity is highly
dependent on cycle-accurate counting
(since the physical floppy disk keeps
spinning whether you write to it or
not). In a normal DOS, this wait
routine (two consecutive JSR and RTS
instructions) always burns exactly 24
cycles. On these disks, it sometimes
burns 16 cycles, sometimes 24. So this
RWTS writes the first 2 sync bytes
faster than usual, then 4 sync bytes at
the usual speed.
Also, the first two times, it writes
$00 instead of $FF. I have no idea what
effect that has on a physical floppy
disk. All valid nibbles have the high
bit set, so maybe it's just dead space?
The entire thing is weird.
This patch should restore the original
wait loop code at $BC69:
T00,S06,$69 change 4CB8B6EAEAEA
                 to 2003BC20C3BC
All the other differences in this RWTS
center around the address and data
prologue and epilogue bytes. The RWTS
apparently only uses a single epilogue
byte ($97), so some of the branch
instructions are subtly different. For
example, checking the address epilogue:
*B98BL
B98B- BD 8C C0
                     LDA $C08C,X
B98E-
       10 FB
                     RPL
                            $R98R
```

```
; usually $DE
B990- Č9 97
                     CMP
                            #$97
                            $B942
B992- D0 AE
                      BNE
B994- EA
B995- BD 8C C0
B998- 10 FB
                     NOP
                     LDA
                           $C08C,X
$B995
                      BPL
; usually $AA
В99A- Č9 00
                     CMP #$00
; usually BNE
B99C- F0 A4
                     BEQ
                            $R942
                     CLC
B99E- 18
B99F- 60
                     RTS
All told, the RWTS needed 15 patches
(including the two I mentioned already,
and the three listed above) to be able
to read and write a standard disk.
A lot of disks need this sort of post-
demuffin patching, and I got tired of
doing it manually, so I wrote a program
to do it for me. It is called,
unsurprisingly, Post-Demuffin Patcher.
It prompts you to select a slot and
drive, then reads the demuffin'd disk,
checks for a modified DOS 3.3-shaped
RWTS, and applies the necessary patches
so the disk can read itself. (It can
also detect and bypass some nibble
checks.) I've included a copy of Post-
Demuffin Patcher on my work disk; the
full source code is currently available
at <http://archive.org/details/
PostDemuffinPatcher4am>.
```

Having already seen several Sunburst disks with identical protection, I knew it would be worth my time to add the detection and repair functions for this crazy RWTS to Post-Demuffin Patcher. So that's exactly what I did. Now, whenever I come across a Sunburst disk (including this one), I can let Post-Demuffin Patcher do the work for me: ES6,D1=demuffin'd copy] JPR#5 JBRUN PDP T00,S03,\$40 change D0 to F0 T00,S03,\$9C change F0 to D0 T00,S06,\$69 change 4CB8B6EAEAEA to 20C3 C20C3BC T00,S08,≸8C change 69BA to A0B9 T00,S03,\$91 change 97 to DE T00,S03,\$9B change 00 to AA change D3 to T00,S03,\$35 DE T00,S03,\$3F change 00 to AΑ T00,S06,\$AE change 97 to DE T00,S06,\$B3 change 00 to AA T00,S06,\$B8 change 00 to EB T00,S02,\$9E change D3 to DE T00,S02,\$A3 change 00 to AΑ T00,S02,\$A8 change 00 to EΒ T00,S02,\$AD change 00 to FF (This is the actual output of the program. Post-Demuffin Patcher prints out the changes it is going to make before it writes them to the disk.)

rest are single-byte patches to get the RWTS to read and write a standard disk with "D5 AA 96"/"D5 AA AD" proloques and "DE AA EB" epilogues. I should point out that Post-Demuffin Patcher is really quite conservative in making patches. It checks a lot of the surrounding code before deciding to patch a specific location (like the RWTS patches). In the case of changing the JMP at \$BE8C from \$BA69 to \$B9Ā0, it actually checks every single byte of code at \$BA69 to ensure that it recognizes the custom routine as the Sunbūrst RWTS. And there were dozens of patches that it didn't make to this disk, because it decided they weren't

The first 2 lines are fixing the nonstandard branch instructions around the second epilogue byte. The next 2 lines are the 2 patches I found earlier. The

needed or it wasn't 100% sure what was going on. JPR#6

And it crashes on boot.

A=FD X=38 Y=1B P=33 S=FB

9D86-

Wait, what?

figured out what the difference was: the original disk uses disk volume 001 instead of the default (254). Why is this a problem? I'm glad you askēd. I used Advanced Demuffin to convert the disk to a standard format, and part of that "standard format" included writing the data to a freshly formatted disk... which uses the default disk volume 254. I could repeat the entire process and format my target disk as disk volume 001, but that will present another problem down the line. Ultimately, I want to distribute this as a .dsk image file, and .dsk image files don't have any way of storing the volume number. They are purely a sector dump; that's just the nature of the file format. The volume number is stored in the address field, but .dsk files don't store the address field. Emulators assume that

every .dsk image uses disk volume 254.

After examining this disk -- which uses identical copy protection to dozens of

other Sunburst disks -- I finally

I'm glad you asked. Besides appearing in every sector's address field on a physical floppy disk, the volume number is stored in three different places when a disk is initialized: \$B7EB (T00,S01,\$EB), in the RWTS parameter table used by boot1 to load DOS from tracks 0-2 ["Beneath Apple DOS", p. 8-351 2. \$AA66 (T01,S09,\$66), in the parsed keyword table used by DOS to load the startup program (and every other file loaded after that) [ibid., p. 8-21**3**  \$B3C1 (T11,S00,\$06), in the UTOC header [ibid., p. 8-32] My (non-working) copy has a \$01 in each of those locations. Because of this mismatch (specifically the first 1, on T00,S01), boot1 fails to load boot2. Literally every sector read fails. Then it jumps to boot2 at \$9D84, which was never loaded, so it crashes. Which is exactly the behavior I saw.

OK, so if .dsk files don't store the volume number, then why does it matter? Using my trusty Disk Fixer sector editor, I changed each of those locations to \$FE.

further copy protection. And because the DOS on my copy now expects disk

volume 254, it also works when I create a .dsk image file out of it and boot it

T00,S01,\$EB change 01 to FE
T01,S09,\$66 change 01 to FE
T11,S00,\$06 change 01 to FE

in an emulator.

(Note to self: add this to a future version of Post-Demuffin Patcher.) Quod erat liberandum.

